2nd International Conference
on Automated Machine Learning:
AutoML 2023 Conference, Potsdam

forester: A Novel Approach to Accessible and Interpretable AutoML for Tree-Based Modeling

Hubert Ruczyński, Anna Kozak
Warsaw University of Technology

forester: A Novel Approach to Accessible and Interpretable AutoML for Tree-Based Modeling

Anna Kozak[1] **Hubert Ruczyński**[1]

[1]Warsaw University of Technology

Abstract The majority of AutoML solutions are developed in Python. However, a large percentage of data scientists are associated with the R language. Unfortunately, there are limited R solutions available with high entry level which means they are not accessible to everyone. To fill this gap, we present the *forester* package, which offers ease of use regardless of the user's proficiency in the area of machine learning.

The *forester* package is an open-source AutoML package implemented in R designed for training high-quality tree-based models on tabular data. It supports regression and binary classification tasks. A single line of code allows the use of unprocessed datasets, informs about potential issues concerning them, and handles feature engineering automatically. Moreover, hyperparameter tuning is performed by Bayesian optimization, which provides high-quality outcomes. The results are later served as a ranked list of models. Finally, the *forester* package offers a vast training report, including the ranked list, a comparison of trained models, and explanations for the best one.

AutoML 2023 Workshop Track

How to build models in R?

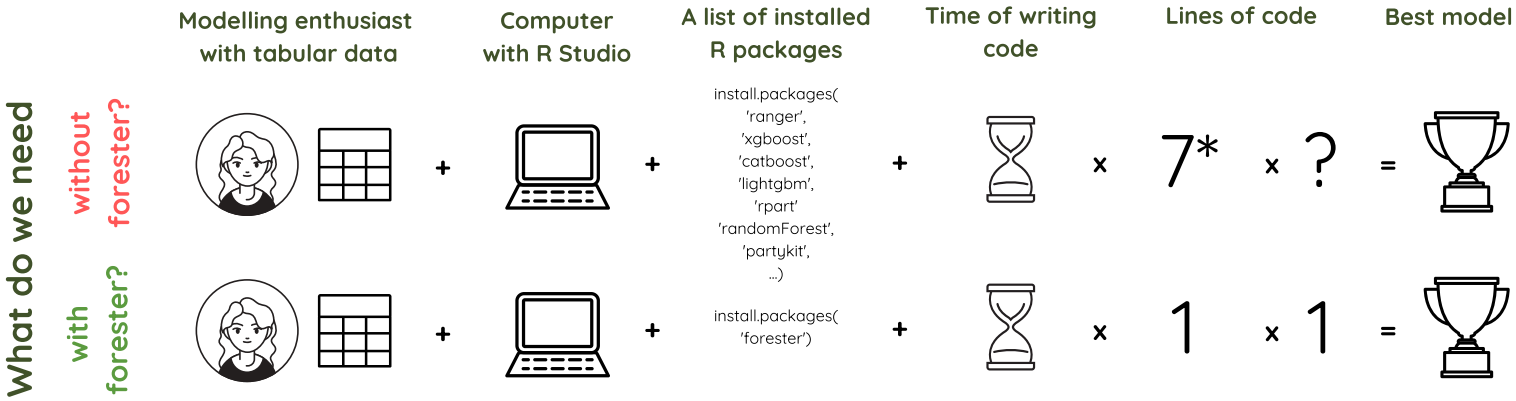

* dependent on the number of packages used

How to use it?

```
library(forester)
data(`lisbon`)
train_output <- train(lisbon,`Price`)
```

What is the forester?

The forester is an AutoML tool in R for **tabular data regression** and **binary classification tasks***, that wraps up all machine learning processes into a single `train()` function, which includes:

1. rendering a brief **data check report**,
2. **preprocessing** the initial dataset enough for models to be trained,
3. **training 5 tree-based models** with default parameters, random search and Bayesian optimization,
4. evaluating them and providing a **ranked list**.

However, that's not everything that the forester has to offer. Via additional functions, the user can easily explain created models with the usage of *DALEX* or generate one of the predefined **reports** including:

1. information about the dataset,
2. in-depth parameters of trained models,
3. visualizations comparing the best models,
4. explanations of the aforementioned models.

forester pipeline

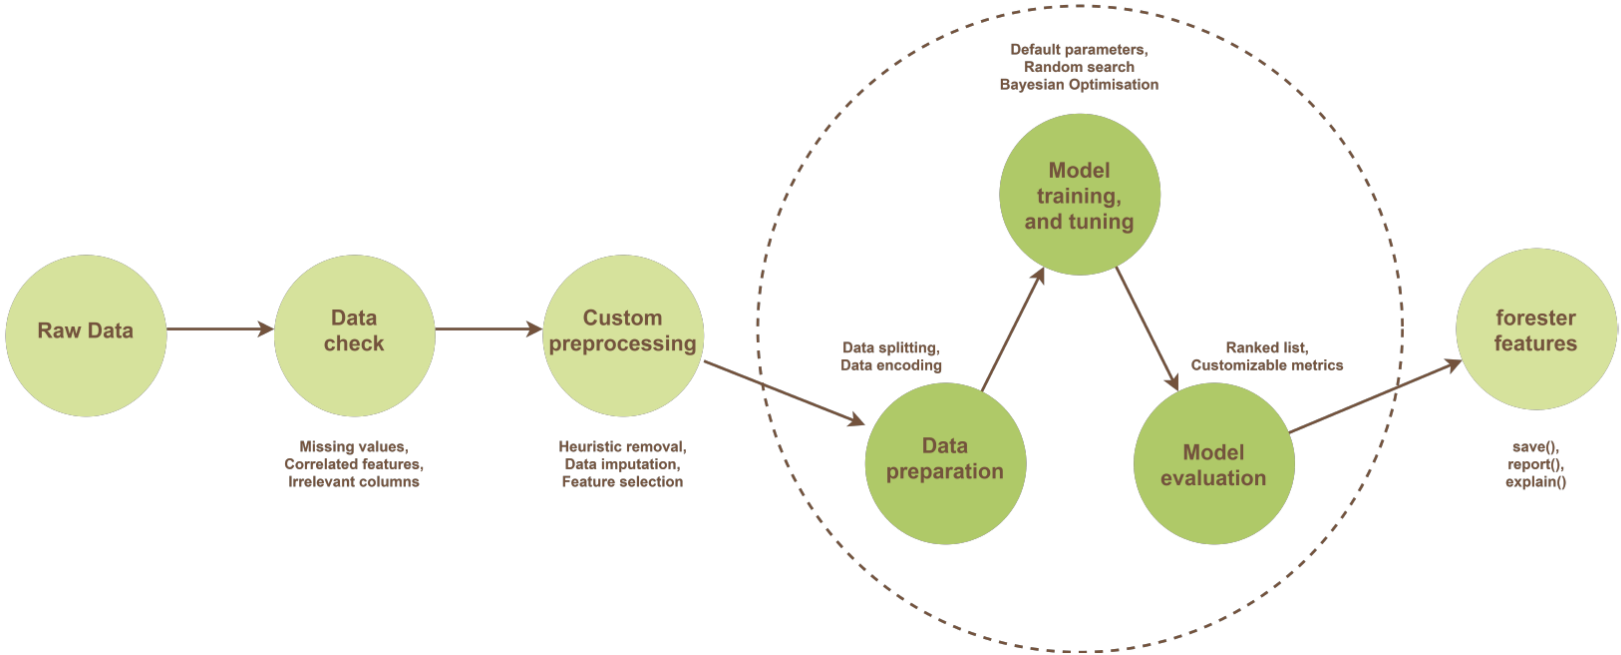

Automatic reports and XAI

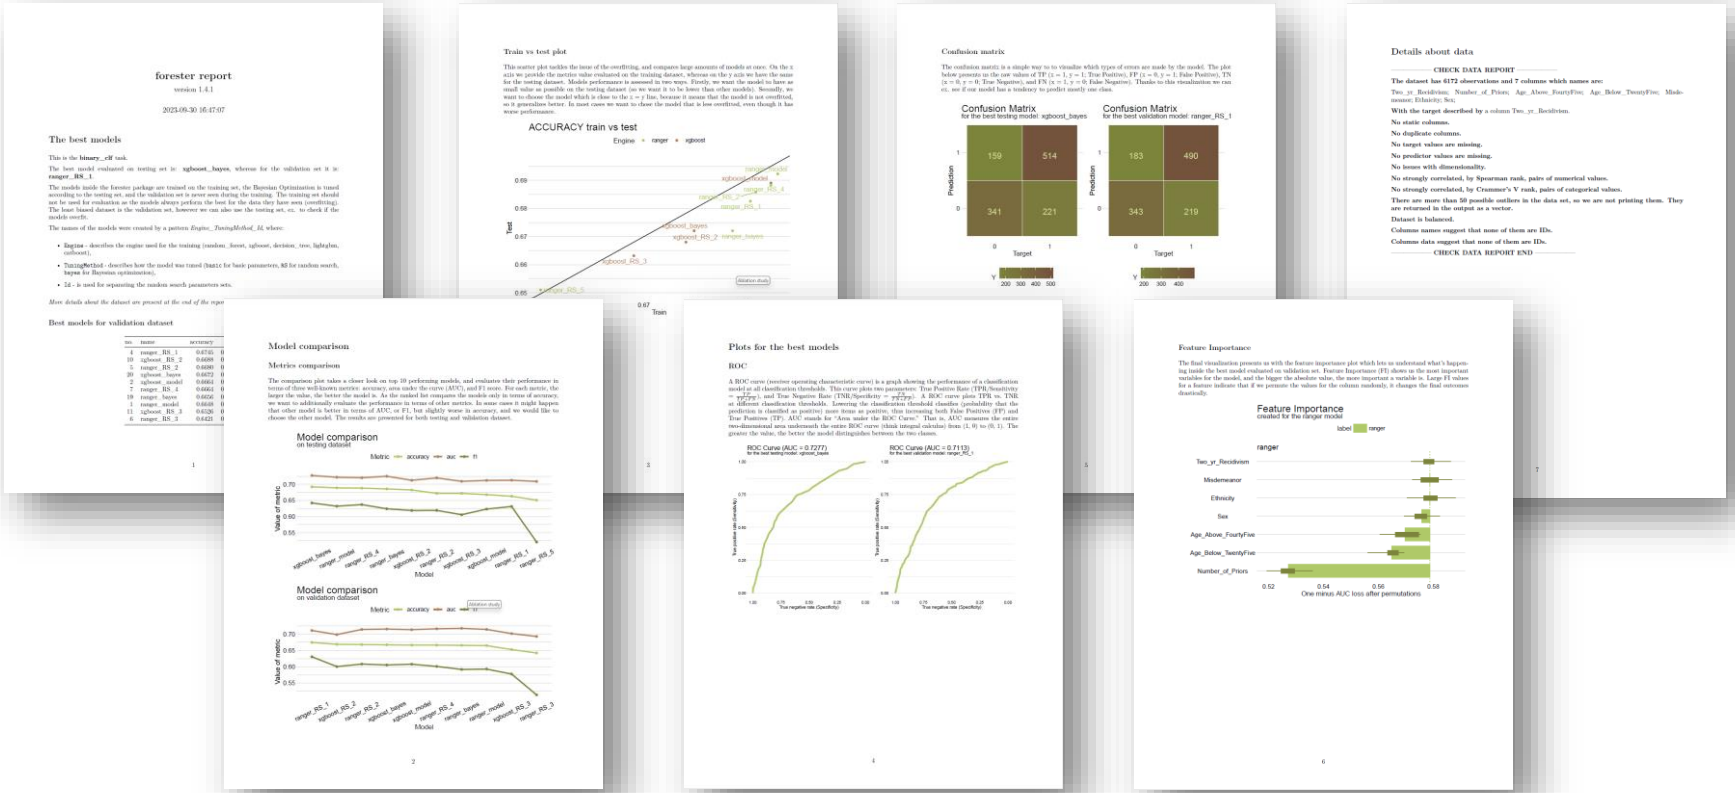

Evaluation

Name	Number of columns	Number of rows
kr-vs-kp	37	3196
breast-w	10	699
credit-approval	16	690
credit-g	21	1000
diabetes	9	768
phoneme	6	5404
banknote-authentication	5	1372
blood-transfusion-service-center	5	748

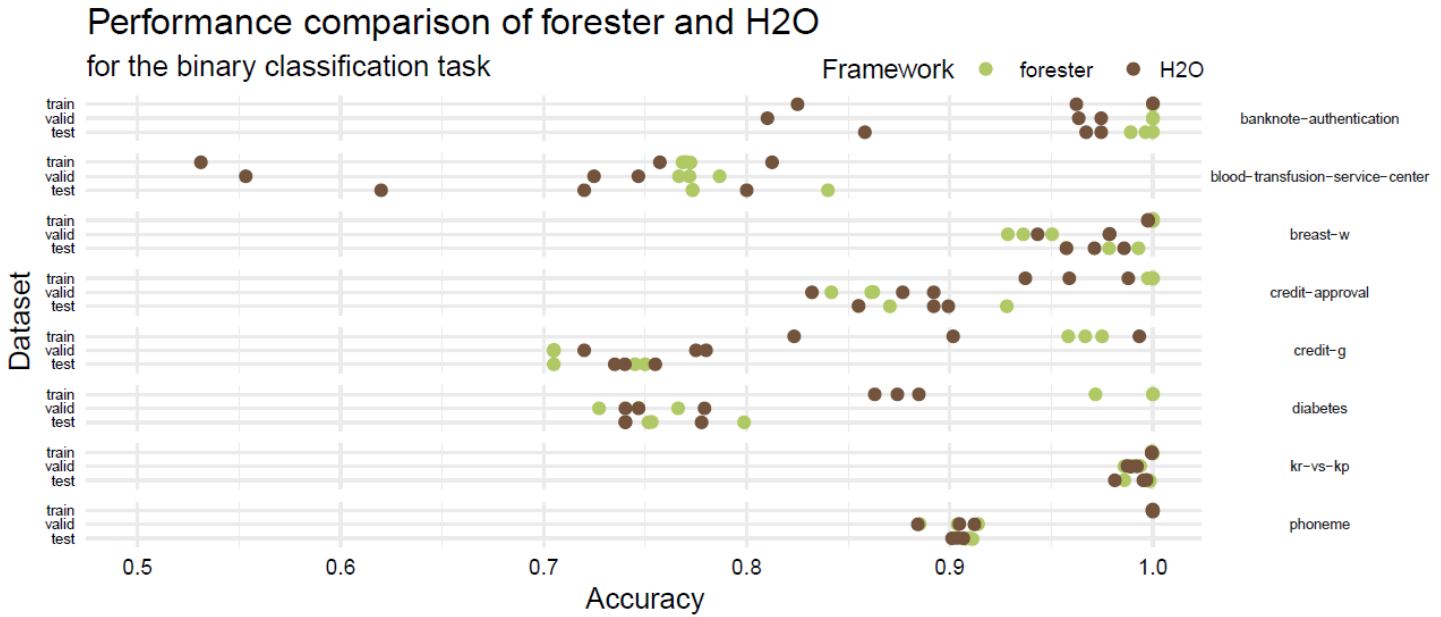

Evaluation

Name	Number of columns	Number of rows
bank32nh	33	8192
wine_quality	12	6497
Mercedes_Benz_Greener_Manufacturing	378	4209
kin8nm	9	8192
pol	49	15000
2dplanes	11	40768
elevators	19	16599

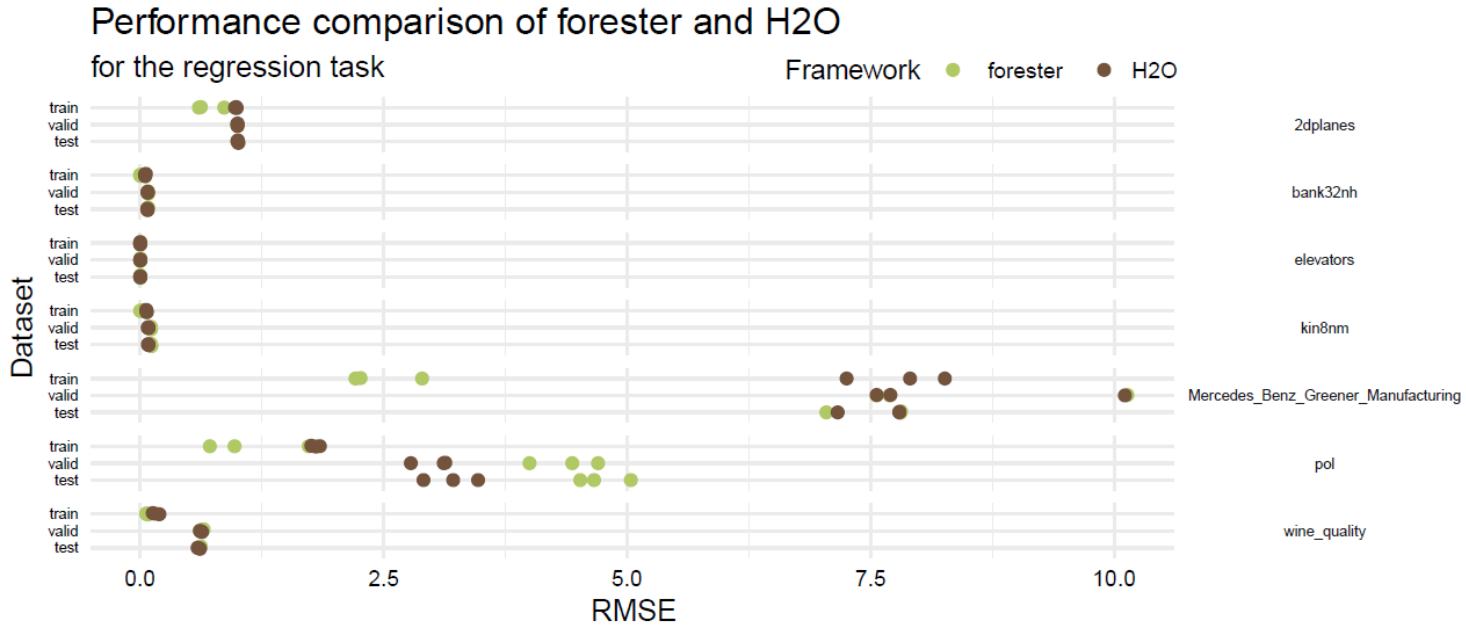

Evaluation

Table 11: The comparison of mean execution times in seconds for the *forester* and *H2O* for binary classification experiments.

task_name	forester	H2O	difference	relative difference
banknote-authentication	818.33	2521.33	-1703	0.28
blood-transfusion-service-center	155.67	555.67	-400	0.26
breast-w	451.33	797.33	-346	0.57
credit-approval	805	1513	-708	0.53
credit-g	2453	4234	-1781	0.58
diabetes	1645.67	2643.67	-998	0.62
kr-vs-kp	451.33	806.67	-355.33	0.57
phoneme	2748.33	3695.33	-947	0.67

Table 12: The comparison of mean execution times in seconds for the *forester* and *H2O* for regression experiments.

task_name	forester	H2O	difference	relative difference
2dplanes	401	1050.67	-649.67	0.38
bank32nh	708.67	1214.67	-506	0.58
elevators	720.33	1435.33	-715	0.5
kin8nm	544.67	1564	-1019.33	0.35
Mercedes_Benz_Greener_Manufacturing	848	1371.67	-523.67	0.61
pol	756	1548.33	-792.33	0.49
wine_quality	1317.33	2130	-812.67	0.63

forester

Public

Edit Pins | Unwatch 11 | Fork 15 | Starred 103

main | 10 branches | 0 tags | Go to file | Add file | Code

HubertR21 Merge pull request #120 from ModelOriented/report ... 090bfe0 last week 226 commits

R	fixes	last week
catboost_info	new dev	last year
data	Version 1.0.0 part 1	last year
docs	update forester page	7 months ago
inst/rmd	fixes	last week
man	fixes	last week
misc	fix	last week
pkgdown	Version 1.0.0 part 2	last year
tests	forester 1.4.1 - enhanced reports	last week
vignettes	knowledge check fixes	10 months ago
.Rbuildignore	new dev	last year
.gitignore	new dev	last year
DESCRIPTION	forester 1.4.1 - enhanced reports	last week
NAMESPACE	report and plots update	last week
NEWS.md	forester 1.4.1 - enhanced reports	last week
README.md	update forester page	7 months ago
forester.Rproj	Version 1.0.0 part 2	last year

README.md

forester: Quick and Simple Tools for Training and Testing of Tree-based Models

A significant amount of time is spent on building models with high performance. Selecting the appropriate model structures, optimizing hyperparameters and explainability are only part of the process of creating a machine learning-based solution. Despite the wide range of structures considered, tree-based models are champions in competitions or

About

Trees are all you need

modeloriented.github.io/forester/

Readme
Activity
103 stars
11 watching
15 forks
Report repository

Releases

No releases published
Create a new release

Packages

No packages published
Publish your first package

Contributors 9

Deployments 74

github-pages last week

+ 73 deployments

Languages

HTML 98.5% R 1.5%

Current directions

The impact of data preparation on the quality of tree-based models created with AutoML forester package

Hubert Ruczyński[1]

[1]Warsaw University of Technology

Abstract Automated Machine Learning (AutoML) solutions are increasingly popular, as they allow data scientists to train high-quality machine learning models with minimal effort. However, the majority of AutoML solutions are developed in Python, leaving R users with limited, and overly complex tools. In this paper, we introduce *forester*, an open-source AutoML package implemented in R that is designed for training high-quality tree-based models on tabular data. The *forester* supports regression, binary classification, and newly implemented survival analysis tasks. The focus put on a single model family, gives us an opportunity to derive conclusions about its behaviour.

Additionally, we introduce a custom preprocessing module, which creates an opportunity to validate a common belief that tree-based models do not require any data preprocessing. We answer this question by conducting a thorough ablation study, of the *forester* package, where we evaluate the impact of dozens of preprocessing strategies. Obtained results let us believe that in the case of the tree-based models family some methods, prove to be more efficient than, others. Finally, we cannot fully agree with the presented thesis, and we provide the reasons supporting our belief.

Ablation study

Data set	Rows	Columns	Static	Duplicate pairs	Missing fields	Dimensional issues	Correlation pairs	Imbalance	ID-like
banknote-authentication	1372	5	0	0	0	No	1	No	No
blood-transfusion-service-center	748	5	0	0	0	No	1	Yes	No
breast-w	699	10	0	0	16	No	9	Yes	No
credit-approval	690	16	0	0	37	No	1	No	No
credit-g	1000	21	0	0	0	No	0	Yes	No
diabetes	768	9	0	0	0	No	0	Yes	No
kr-vs-kp	3196	37	4	0	0	Yes	0	No	No
phoneme	5403	6	0	0	0	No	0	Yes	No

Data set	Rows	Columns	Static	Duplicate pairs	Missing fields	Dimensional issues	Correlation pairs	Imbalance	ID-like
2dplanes	40768	11	0	0	0	No	0	No	No
bank32nh	8192	33	0	0	0	Yes	0	Yes	No
elevators	16599	19	2	0	0	No	11	Yes	No
kin8nm	8192	9	0	0	0	No	0	No	No
Mercedes_Benz_ Greener_Manufacturing	4209	378	145	134	0	Yes	522	No	Yes
pol	15000	49	22	156	0	Yes	2	Yes	No
wine_quality	6497	12	0	0	0	No	1	Yes	No

Ablation study

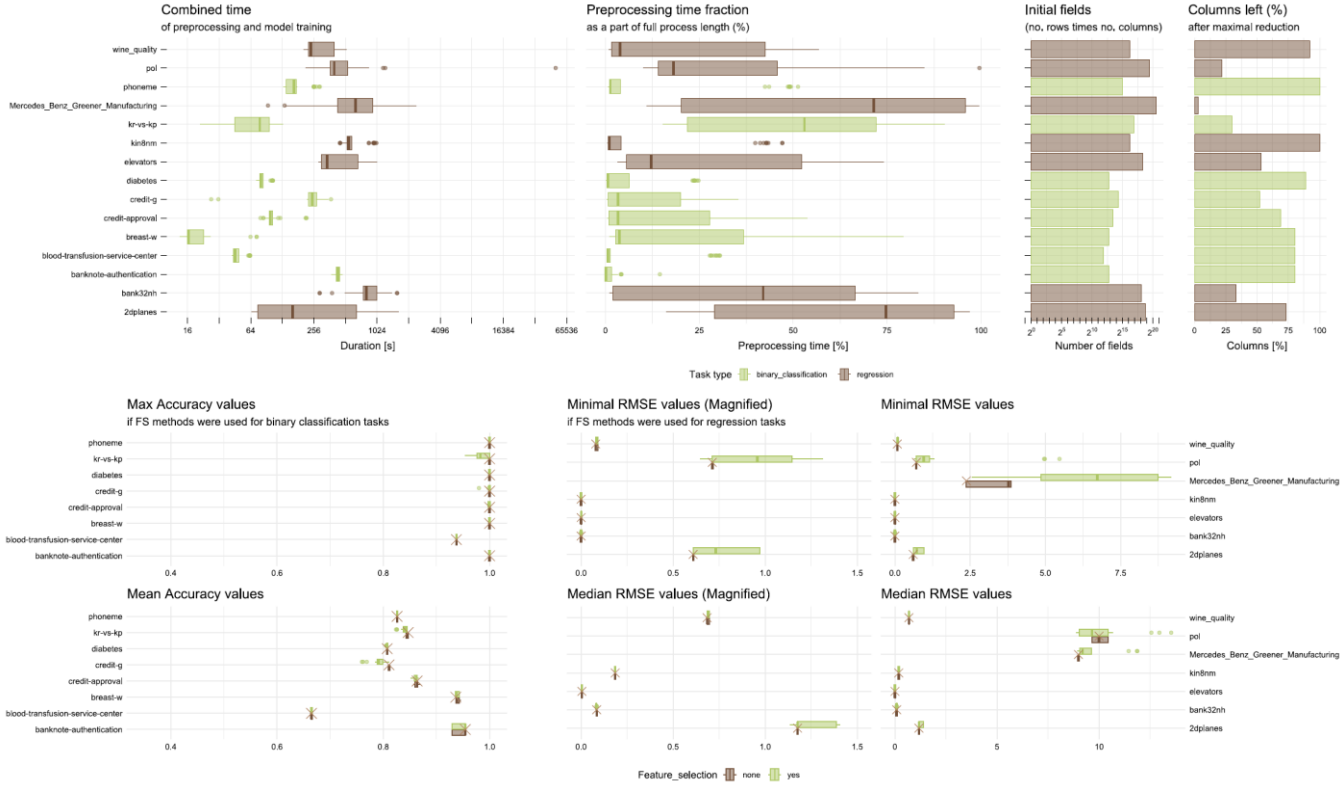

Master's degree

1. Implementation of multiclass classification (MC).
2. Adding a few more datasets for current tasks.
3. Conducting preprocessing methods ablations study for MC.
4. Evaluation of various preprocessing strategies on new data.
5. Uploading forester to CRAN Repository.


Thank you for attention!

forester: an R package for automated building of tree-based machine learning models

Anna Kozak, Adrianna Grudzień, Hubert Ruczyński, Patryk Słowakiewicz
MI2.AI Group, Faculty of Mathematics and Information Science, Warsaw University of Technology

Introduction

A significant amount of time is spent on building models with high performance. Selecting the appropriate model structures, optimising hyperparameters and explainability are only part of the process of creating a machine learning-based solution. Despite the wide range of structures considered, tree-based models are champions in competitions or hackathons. So, aren't tree-based models enough? They are, and that's why **we want to fully automate the process of training tree-based models so that even the newcomers can easily build**, train and understand these powerful prediction tools. At the same time, **the experienced users gain a powerful tool for making high-quality baseline models** for new tasks, they start working with.

What is the *forester*?

The *forester* is **an autoML tool in R** that wraps up all machine learning processes into a single `train()` function, which includes:

1. rendering a brief data check report,
2. preprocessing initial dataset enough for models to be trained,
3. training 5 tree-based models with default parameters, random search and Bayesian optimisation,
4. evaluating them and providing a ranked list.

However, that's not everything that the *forester* has to offer. Via additional functions, the user can easily explain created models with the usage of *DALEX* or generate one of the predefined reports including:

1. information about the dataset,
2. in-depth parameters of trained models,
3. visualisations comparing the best models,
4. explanations of the aforementioned models.

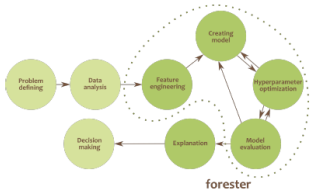

Why tree-based models?

Tree-based models, especially **XGBoost are extremely popular amongst winners in Kaggle competitions** and they firmly show their superiority with tabular data, not only in terms of fast computations. Moreover, the researchers also prove that **tree-based models are superior to deep learning neural networks** because they don't suffer from uninformative columns presence and are not biased toward overly smoothed solutions.

MI RESEARCH

Package structure

With functions in *forester* package users can create a well-tuned tree-based model with a unified, simple formula. With the usage of only two required parameters: the raw, not preprocessed dataset and target column name, the user is able to achieve satisfying results. The *forester* automatically handles the "ugly" part for you.

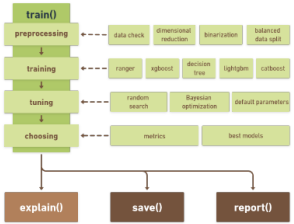

For whom is this package created?

The *forester* is designed for beginners in data science, but also for more experienced users. They get an easy-to-use tool that can be used to prepare high-quality baseline models for comparison with more advanced methods or a set of output parameters for more thorough optimisations. **Tree-based models are created in just one line of code.** The package differentiates itself in this aspect from powerful autoML frameworks like *mlr3* and *H2O*.

	forester	mlr3	H2O
easy to use	✓		
preprocesing	✓	✓	
autoML	✓	✓	✓
feature selection	◔	✓	✓
model tuning	✓	✓	✓
vizualization	✓		✓
explanation	✓		✓
report	✓		✓

Contact info

anna.kozak@pw.edu.pl
grudziena@outlook.com
hruczynski21@interia.pl
slowakiewiczpatryk@outlook.com
https://github.com/ModelOriented/forester

forester: growing transparent tree-based models for everyone

Anna Kozak[1], Adrianna Grudzień[1], Hubert Ruczyński[1],
Patryk Słowakiewicz[1], Przemysław Biecek[1,2]
[1]MI2.AI, Warsaw University of Technology [2]MI2.AI, University of Warsaw

Let's talk about AutoML, tree-based models, explainable AI (XAI), exploratory data analysis (EDA)!

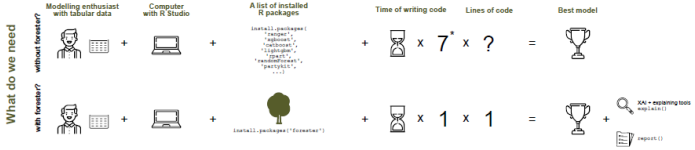

How to build tree-based models in R?

What is *forester*?

- full automation of the process of training tree-based models
- no demand for ML expertise
- powerful tool for making high-quality baseline models for experienced users

The *forester* package is **an AutoML tool in R** that wraps up all machine learning processes into a single `train()` function, which includes:

1. rendering a brief **data check** report,
2. **preprocessing** initial dataset enough for models to be trained,
3. **training** 5 tree-based models with default parameters, random search and Bayesian optimisation,
4. **evaluating** them and providing a ranked list.

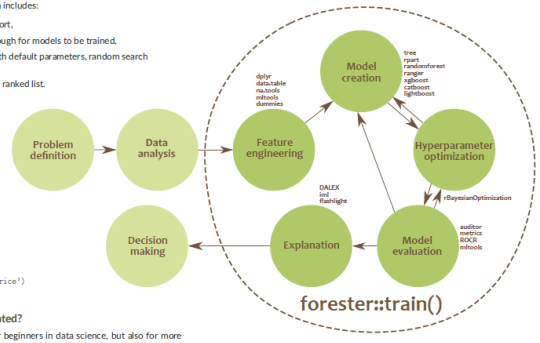

How to use it?

```
library(forester)
data('liabon')
train_output <- train(liabon, 'Price')
```

For whom is this package created?

The *forester* package is designed for beginners in data science, but also for more experienced users. They get an easy-to-use tool that can be used to prepare high-quality baseline models for comparison with more advanced methods or a set of output parameters for more thorough optimisations.

Contact info

anna.kozak@pw.edu.pl
https://github.com/ModelOriented/forester

References

P. Biecek. DALEX: Explainers for Complex Predictive Models in R. *Journal of Machine Learning Research*, 19(84):1–5, 2018. URL https://jmlr.org/papers/v19/18-416.html.
A. Kozak, H. Ruczyński, P. Słowakiewicz, A. Grudzień, and P. Biecek. *forester: Quick and Simple Tools for Training and Testing of Tree-based Models*, 2022. URL https://github.com/ModelOriented/forester. R package version 1.0.0.

Prepare meaningful report less than in 60 seconds!

As data scientists, we are fully aware that there are some time expensive processes in out work. One of them is creating a report with meaningful results. That's why one of the most powerful *forester* feature, which makes it a efficient tool for both experienced users and the newcomers, is a `report()` function. This single-line command is designed to **provide a holistic view on the outcomes of the ML process** happening inside of the *forester*.

See the report yourself!

forester: A Novel Approach to Accessible and Interpretable AutoML for Tree-Based Modeling

Anna Kozak [1], Hubert Ruczyński[1]

[1]Warsaw University of Technology

Let's talk about AutoML, tree-based models, explainable AI (XAI), and exploratory data analysis (EDA)!

How to build tree-based models in R?

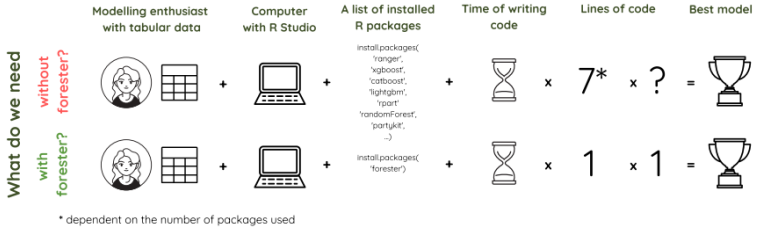

* dependent on the number of packages used

What is *forester*?

- Full automation of the process of training tree-based models,
- No demand for ML expertise,
- Powerful tool for making high-quality baseline models for experienced users.

The *forester* package is **an AutoML tool in R** that wraps up all machine learning processes
into a single `train()` function, which includes:

1. Rendering a brief **data check** report,
2. **Preprocessing** initial dataset enough for models to be trained,
3. **Training** 5 tree-based models with default parameters, random search and Bayesian optimisation,
4. **Evaluating** them and providing a ranked list.

For whom is this package created?

The *forester* package is designed for beginners in data science, but also for more experienced users. They get an easy-to-use tool that can be used to prepare high-quality baseline models for comparison with more advanced methods or a set of output parameters for more thorough optimisations.

How to use it?

```
library(forester)
data('lisbon')
train_output <- train(lisbon, 'Price')
```

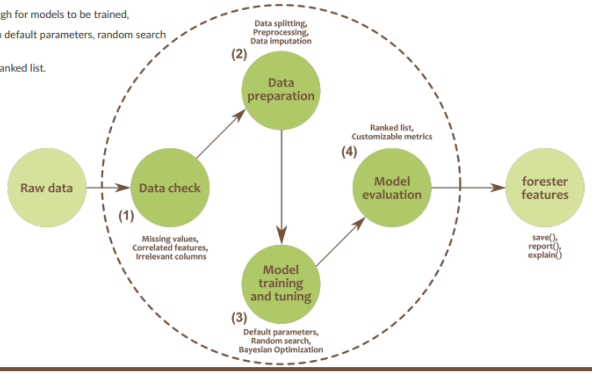

Prepare meaningful report less than in 60 seconds!

As data scientists, we are fully aware that there are some time expensive processes in out work. One of them is creating a report with meaningful results. That's why one of the most powerful *forester* feature, which makes it a efficient tool for both experienced users and the newcomers, is a `report()` function. This single-line command is designed to **provide a holistic view on the outcomes of the ML process** happening inside of the *forester*.

Simple doesn't mean worse!

According to our experiments, the forester package achieves competitive results in much shorter time in comparison to well-known H2O AutoML tool. We have compared their performance on 8 binary classification, and 7 regression tasks, and the calculations were repeated 3 times for each dataset and framework. The forester outperformed H2O most of the times, even though the latter package's training lasted 2 times longer on average.

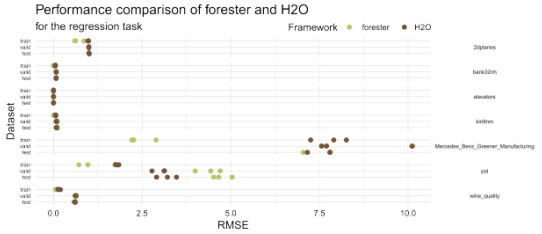

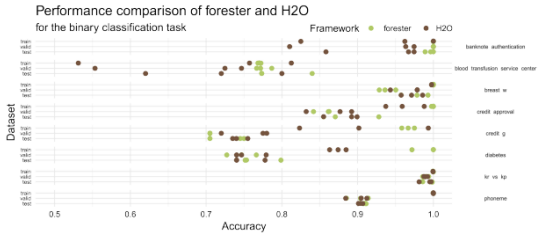

Contact info

anna.kozak@pw.edu.pl
hruczynski21@interia.pl
https://github.com/ModelOriented/forester

Paper

GitHub

References

L. Grinsztajn, E. Oyallon, and G. Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2022. URL https://openreview.net/forum?id=Fp7__phQszn.

A. Kozak, H. Ruczyński, P. Słowakiewicz, A. Grudzień, and P. Biecek. *forester: Quick and Simple Tools for Training and Testing of Tree-based Models*, 2023. URL https://github.com/ModelOriented/forester. R package version 1.1.4.

Investigating the Efficiency of Tree-based Models for Tabular Data with *forester* Package

Hubert Ruczyński, Anna Kozak

Warsaw University of Technology

What is *forester*?

The *forester* package is **an AutoML tool in R** that wraps up all machine learning processes into a single `train()` function, which includes:

1. rendering a brief **data check** report,
2. **preprocessing** initial dataset enough for models to be trained,
3. **training** 5 tree-based models with default parameters, random search and Bayesian optimization,
4. **evaluating** them and providing a ranked list.

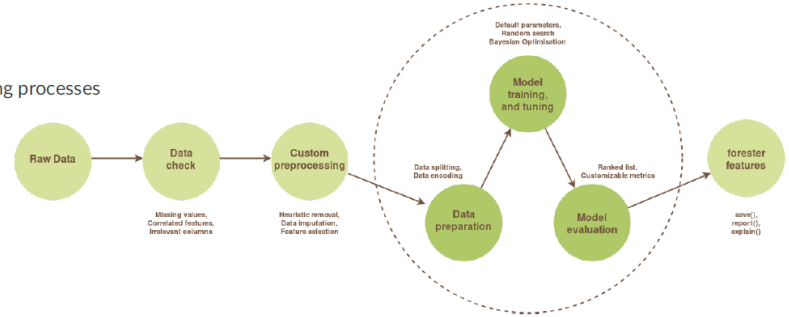

Simple doesn't mean worse!

According to our experiments, the forester package achieves competitive results in much shorter time in comparison to well-known H2O AutoML tool. We have compared their performance on 8 binary classification, and 7 regression tasks, and the calculations were repeated 3 times for each dataset and framework. The forester outperformed H2O most of the times, even though the latter package's training lasted 2 times longer on average.

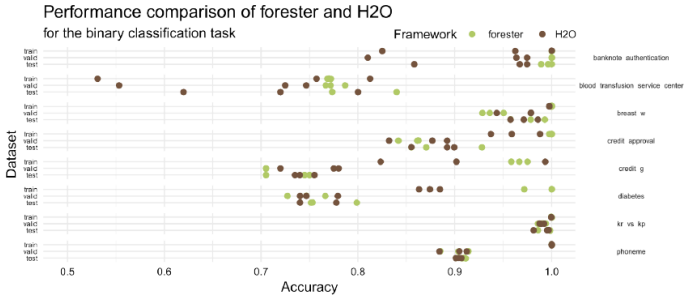

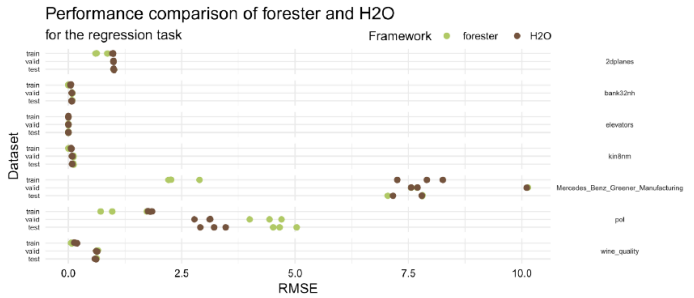

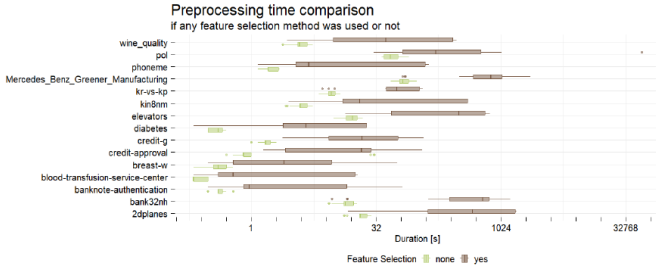

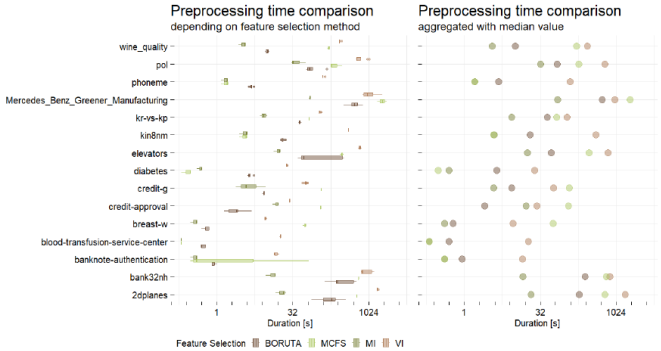

Time complexity

The most important findings consider feature selection (FS) methods. They are the most expensive part of preprocessing, and their execution times differ significantly between the methods. The right plot shows us that Mutual Information (MI) based selection method, and BORUTA are relatively fast, whereas Monte Carlo Feature Selection (MCFS), and Variable Importance (VI) are rather slow.

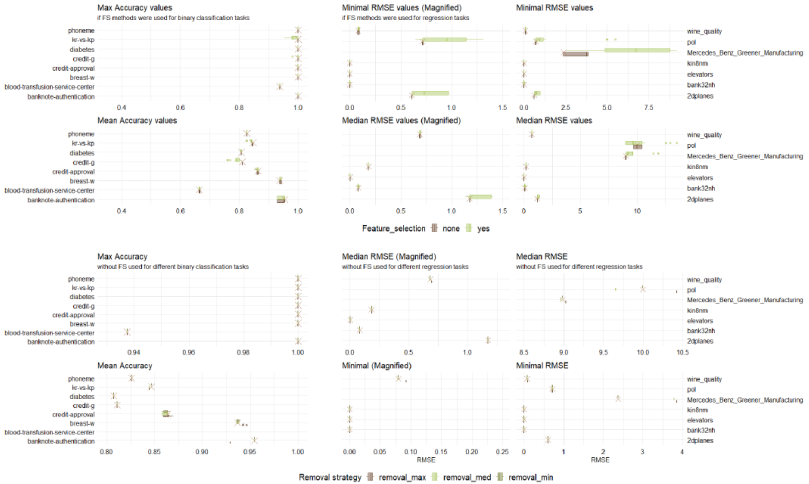

Feature selection impact on performance

FS methods are responsible for unstable results, and in most cases, its usage leads to worse results than for baseline methods marked with **X**. In some cases however, with FS methods we can obtain better results.
When we consider preprocessing strategies based on heuristic removals, the results are less significant, but in most cases lead to enhancements of the results.

Contact info

hruczynski21@interia.pl
https://github.com/ModelOriented/forester

References

A. Kozak and H. Ruczyński. forester: A Novel Approach to Accessible and Interpretable AutoML for Tree-Based Modeling. In *AutoML Conference 2023 (ABCD Track)*, 2023. URL https://openreview.net/forum?id=Q3DWpGoX7PD.